



**Homework #2**

**(HW2 is Due on Friday, September 21 11.59pm)**

# Adil Ashish Kumar

(put your full name above (incl. any nicknames); 5 points)

Note: This is an individual homework. Discussing this homework with your classmates is a violation of the Honor Code. If you borrow code from somewhere else, please add a comment in your code to make it clear what the source of the code is (e.g., a URL would sufficient). If you borrow code and you don't provide the source, it is a violation of the Honor Code.

Total grade: \_\_\_\_\_ out of \_\_\_\_150\_\_\_\_ points

**1) (15 points) Would you frame the problem of e-mail spam detection as a supervised learning problem or an unsupervised learning problem? Please justify your answer.**

The problem is a supervised learning problem because the goal is to classify an email as spam or not. The target variable in this case would be spam/no spam column having 1 for spam and 0 otherwise. Here, the assumption is that we have a dataset with a column to label an email as spam or not. Unlike in unsupervised learning, there is a specific goal/purpose (spam email prediction) in this problem, which makes it a supervised learning problem. This would be a classification problem since the target variable would be categorical.

**2) (15 points) What is a test set and why would you want to use it?**

A test set is a part of the data on which we do not train the model. After the model is trained on training data, we evaluate model performance by making predictions on the test dataset and compare the predictions with the target variable in test. Generally, the about 20-30% of the data is set aside for test and about 70-80% of the data would be used to train the model. These numbers may vary based on the size of the dataset. The test dataset should not be used as part of the modeling process, but rather to evaluate the performance of a model.

A test dataset is crucial to generalization performance. Without testing the model on data it has not seen, it is hard to determine the model accuracy. Metrics such as out of sample accuracy, Precision, recall, F measure and error rate give a picture of model performance on test data. With help of test data, we can also create fitting graphs to identify the problem of overfitting by observing error% as model complexity is increased.

**3) (20 points) What are the similarities and differences of decision trees and logistic regression? When might you prefer to use one over another?**

**Differences:**

Tree models are more flexible than logistic regression since they can handle nonlinear relationships in the data. In classification trees, decision boundaries are perpendicular to the instance space axes and there can be multiple decision boundaries, but in logistic regression decision boundaries can have any orientation. In Logistic regression, a single decision surface is placed through the entire space, while in tree models, the instance space can be split into small regions. Tree models consider 1 attribute at a time for each split from parent to child node, but logistic regression considers all predictor variables together in model. Model interpretability is better for decision trees since the output of model is a set of rules that is easy to understand, while output of logistic regression is a set of model coefficients which are not easy to understand for non-statistically savvy audience.

**Similarities:** Both models can be used for classification problems, but logistic regression cannot be used for regression. Both models can overcome overfitting problems – tree models use pre and post pruning techniques while logistic regression uses regularization and feature selection. Probability estimation can be done by both models.

**Usage:** Logistic regression can be used only for classification problems but trees can be used for both classification and regression problems. Logistic regression is preferred over tree models for smaller datasets. Tree models need more data in general. Tree models can handle datasets with non linear relationships, missing values, better than logistic regression. Tree models have better model comprehensibility than logistic regression. Flexibility of tree models give it an advantage for use on larger datasets.

**4) (25 points)** You have a fraud detection task (predicting whether a given credit card transaction is “fraud” vs. “non-fraud”) and you built a classification model for this purpose. For any credit card transaction, your model estimates the probability that this transaction is “fraud”. The following table represents the probabilities that your model estimated for the validation dataset containing 10 records.

Actual Class (from validation data)	Estimated Probability of Record Belonging to Class “fraud”
fraud	0.95
fraud	0.91
fraud	0.75
non-fraud	0.67
fraud	0.61
non-fraud	0.46
fraud	0.42
non-fraud	0.25
non-fraud	0.09
non-fraud	0.04

Based on the above information, answer the following questions:

- a) What is the overall accuracy of your model, if the chosen probability cutoff value is 0.3? What is the overall accuracy of your model, if the chosen probability cutoff value is 0.8?

Actual Class (from validation data)	Estimated Probability of Record Belonging to Class “fraud”	Prediction (0.3)	Prediction (0.8)	Prediction (0.7)	Prediction (0.4)
fraud	0.95	Fraud	Fraud	Fraud	Fraud
fraud	0.91	Fraud	Fraud	Fraud	Fraud
fraud	0.75	Fraud	non-fraud	Fraud	Fraud
non-fraud	0.67	Fraud	non-fraud	Non	Fraud
fraud	0.61	Fraud	non-fraud	Non	Fraud
non-fraud	0.46	Fraud	non-fraud	Non	Fraud
fraud	0.42	Fraud	non-fraud	Non	Fraud
non-fraud	0.25	non-fraud	non-fraud	Non	Non
non-fraud	0.09	non-fraud	non-fraud	Non	Non
non-fraud	0.04	non-fraud	non-fraud	Non	Non

**Cutoff = 0.3, accuracy =  $8/10 = 80\%$**

**Cutoff = 0.8, accuracy =  $7/10 = 70\%$**

- b) What probability cutoff value should you choose, in order to have Precision fraud = 100% for your model? (Explain.) What is the overall accuracy of your model in this case?**

A cutoff value of 0.7 would result in a 100% precision fraud. The overall accuracy for the model would be 80%  
Below is the confusion matrix for the same:

	Actual fraud	Actual Non fraud
Predicted fraud	3	0
Predicted Non fraud	2	5

- c) What probability cutoff value should you choose, in order to have Recall fraud = 100% for your model? (Explain.) What is the overall accuracy of your model in this case?**

A cutoff value of 0.4 would result in a 100% recall fraud. The overall accuracy for the model would be 80%  
Below is the confusion matrix for the same:

	Actual fraud	Actual Non fraud
Predicted fraud	5	2
Predicted Non fraud	0	3

- d) Draw an ROC curve for your model.**

5) (70 points) [Mining publicly available data. Please implement the following models both with Rapidminer and Python but explore the data (e.g., descriptive statistics etc.) just with Python]

Please use the dataset on breast cancer research from this link: <http://archive.ics.uci.edu/ml/machine-learning-databases/breast-cancer-wisconsin/wdbc.data> [Note: Rapidminer can import .data files in the same way it can import .csv files. For Python please read the data directly from the URL without downloading the file on your local disk.] The description of the data and attributes can be found at this link: <http://archive.ics.uci.edu/ml/machine-learning-databases/breast-cancer-wisconsin/wdbc.names> . Each record of the data set represents a different case of breast cancer. Each case is described with 30 real-valued attributes: attribute 1 represents case id, attributes 3-32 represent various physiological characteristics, and attribute 2 represents the type (benign or malignant). If the dataset has records with missing values, you can filter out these records using Python and/or Rapidminer before proceeding with the models. Alternatively, if the data set has missing values, you could infer the missing values.

Perform a predictive modeling analysis on this same dataset using the a) k-NN technique (for k=3) and b) Logistic Regression. Please be specific about what other parameters you specified for your models.

Present a brief overview of your predictive modeling process, explorations, and discuss your results.

Compare the k-NN model with the Logistic Regression model: which performs better? Make sure you present information about the model “goodness” (i.e., confusion matrix, predictive accuracy, precision, recall, f-measure). Please be clear about any assumptions you might make when you choose the best performing model.

Please show screenshots of the models you have built with Rapidminer and Python, show screenshots of the performance results, and the parameters you have specified.

Brief overview of modeling process:

### 1) Data Input and Preliminary analysis

```
# To write a Python 2/3 compatible codebase, the first step is to add this line to the top of each module
from __future__ import division, print_function, unicode_literals

import numpy as np # np is an alias pointing to numpy
import pandas as pd # pd is an alias pointing to pandas
pd.set_option('display.max_columns', 50) #increasing no columns to display
pd.set_option('display.width', 130) #increasing panda output window width
# reading data from URL
df = pd.read_csv("http://archive.ics.uci.edu/ml/machine-learning-databases/breast-cancer-wisconsin/wdbc.data", header=None)
df.head() # understanding the data by viewing first 5 rows
print(df.describe(include='all')) # simple statistics on each column in data
print()
print(df.shape) # no of rows and columns in data
print()
print(df.isnull().any()) # to check if any column has missing values
```

In this step, the data (.data file) is read in python directly from the URL, with “header = None” to indicate the first row is not a header row. The output width and column limit is adjusted to make sure we can explore the data properly since the data has 32 columns. A quick preliminary analysis is done on the data to get an idea about the structure of the data. The df.shape command tells us that the data has 569 rows and 32 columns. The df.isnull().any() command shows there are no missing values in the data, which indicates no missing value treatment needs to be carried out.

Using df.head(), a sample of the data is printed to get an idea of the structure of the data. As given by the data description, the first column is the case ID, and each row of the data is a unique case. The second column is the target variable, indicating malignancy or benign. All the other columns are numeric attributes.

```

(569, 32)
0 False
1 False
2 False
3 False
4 False
5 False
6 False
7 False
8 False
9 False
10 False
11 False
12 False
13 False
14 False
15 False
16 False
17 False
18 False
19 False
20 False
21 False
22 False
23 False
24 False
25 False
26 False
27 False
28 False
29 False
30 False
31 False
dtype: bool

      0 1 2 3 4 5 6 7 8 9 10 11 12 13 \
14 False 0 842302 M 17.99 10.38 122.80 1001.0 0.11840 0.27760 0.3001 0.14710 0.2419 0.07871 1.0950 0.9053
15 False 1 842517 M 20.57 17.77 132.90 1326.0 0.08474 0.07864 0.0869 0.07017 0.1812 0.05667 0.5435 0.7339
16 False 2 84300903 M 19.69 21.25 130.00 1203.0 0.10960 0.15990 0.1974 0.12790 0.2069 0.05999 0.7456 0.7869
17 False 3 84348301 M 11.42 20.38 77.58 386.1 0.14250 0.28390 0.2414 0.10520 0.2597 0.09744 0.4956 1.1560
18 False 4 84358402 M 20.29 14.34 135.10 1297.0 0.10030 0.13280 0.1980 0.10430 0.1809 0.05883 0.7572 0.7813
19 False
20 False
21 False
22 False
23 False
24 False
25 False
26 False
27 False
28 False
29 False
30 False
31 False
dtype: bool

      14 15 16 17 18 19 20 21 22 23 24 25 26 \
21 False 0 8.589 153.40 0.006399 0.04904 0.05373 0.01587 0.03003 0.006193 25.38 17.33 184.60 2019.0 0.1622
22 False 1 3.398 74.08 0.005225 0.01308 0.01860 0.01340 0.01389 0.003532 24.99 23.41 158.80 1956.0 0.1238
23 False 2 4.585 94.03 0.006150 0.04006 0.03832 0.02058 0.02250 0.004571 23.57 25.53 152.50 1709.0 0.1444
24 False 3 3.445 27.23 0.009110 0.07458 0.05661 0.01867 0.05963 0.009208 14.91 26.50 98.87 567.7 0.2098
25 False 4 5.438 94.44 0.011490 0.02461 0.05688 0.01885 0.01756 0.005115 22.54 16.67 152.20 1575.0 0.1374
26 False
27 False
28 False
29 False
30 False
31 False
dtype: bool

      27 28 29 30 31
0 0.6656 0.7119 0.2654 0.4601 0.11890
1 0.1866 0.2416 0.1860 0.2750 0.08902
2 0.4245 0.4504 0.2430 0.3613 0.08758
3 0.8663 0.6869 0.2575 0.6638 0.17300
4 0.2050 0.4000 0.1625 0.2364 0.07678

      0 1 2 3 4 5 6 7 8 9 \
count 5.690000e+02 569 569.000000 569.000000 569.000000 569.000000 569.000000 569.000000 569.000000 569.000000
unique NaN 2 NaN NaN NaN NaN NaN NaN NaN NaN
top NaN B NaN NaN NaN NaN NaN NaN NaN NaN NaN
freq NaN 357 NaN NaN NaN NaN NaN NaN NaN NaN NaN
mean 3.037183e+07 NaN 14.127292 19.289649 91.969033 654.889104 0.096360 0.104341 0.088799 0.048919
std 1.250206e+08 NaN 3.524049 4.301036 24.298981 351.914129 0.014064 0.052813 0.079720 0.038803
min 8.670000e+03 NaN 6.981000 9.710000 43.790000 143.500000 0.052630 0.019380 0.000000 0.000000
25% 8.692180e+05 NaN 11.700000 16.170000 75.170000 420.300000 0.086370 0.064920 0.029560 0.020310
50% 9.060240e+05 NaN 13.370000 18.840000 86.240000 551.100000 0.095870 0.092630 0.061540 0.033500
75% 8.813129e+06 NaN 15.780000 21.800000 104.100000 782.700000 0.105300 0.130400 0.130700 0.074000
max 9.113205e+08 NaN 28.110000 39.280000 188.500000 2501.000000 0.163400 0.345400 0.426800 0.201200

      10 11 12 13 14 15 16 17 18 19
\
count 569.000000 569.000000 569.000000 569.000000 569.000000 569.000000 569.000000 569.000000 569.000000 569.000000
unique NaN NaN NaN NaN NaN NaN NaN NaN NaN NaN
top NaN NaN NaN NaN NaN NaN NaN NaN NaN NaN
freq NaN NaN NaN NaN NaN NaN NaN NaN NaN NaN
mean 0.181162 0.062798 0.405172 1.216853 2.866059 40.337079 0.007041 0.025478 0.031894 0.011796
std 0.027414 0.007060 0.277313 0.551648 2.021855 45.491006 0.003003 0.017908 0.030186 0.006170
min 0.106000 0.049960 0.111500 0.360200 0.757000 6.802000 0.001713 0.002252 0.000000 0.000000
25% 0.161900 0.057700 0.232400 0.833900 1.606000 17.850000 0.005169 0.013080 0.015090 0.007638
50% 0.179200 0.061540 0.324200 1.108000 2.287000 24.530000 0.006380 0.020450 0.025890 0.010930
75% 0.195700 0.066120 0.478900 1.474000 3.357000 45.190000 0.008146 0.032450 0.042050 0.014710
max 0.304000 0.097440 2.873000 4.885000 21.980000 542.200000 0.031130 0.135400 0.396000 0.052790

      20 21 22 23 24 25 26 27 28 29
\
count 569.000000 569.000000 569.000000 569.000000 569.000000 569.000000 569.000000 569.000000 569.000000 569.000000
unique NaN NaN NaN NaN NaN NaN NaN NaN NaN NaN
top NaN NaN NaN NaN NaN NaN NaN NaN NaN NaN
freq NaN NaN NaN NaN NaN NaN NaN NaN NaN NaN
mean 0.020542 0.003795 16.269190 25.677223 107.261213 880.583128 0.132369 0.254265 0.272188 0.114606
std 0.008266 0.002646 4.833242 6.146258 33.602542 569.356993 0.022832 0.157336 0.208624 0.065732
min 0.007882 0.000895 7.930000 12.020000 50.410000 185.200000 0.071170 0.027290 0.000000 0.000000
25% 0.015160 0.002248 13.010000 21.080000 84.110000 515.300000 0.116600 0.147200 0.114500 0.064930
50% 0.018730 0.003187 14.970000 25.410000 97.660000 686.500000 0.131300 0.211900 0.226700 0.099930
75% 0.023480 0.004558 18.790000 29.720000 125.400000 1084.000000 0.146000 0.339100 0.382900 0.161400
max 0.078950 0.029840 36.040000 49.540000 251.200000 4254.000000 0.222600 1.058000 1.252000 0.291000

      30 31
count 569.000000 569.000000
unique NaN NaN
top NaN NaN
freq NaN NaN
mean 0.290076 0.083946
std 0.061867 0.018061
min 0.156500 0.055040
25% 0.250400 0.071460
50% 0.282200 0.080040
75% 0.317900 0.092080
max 0.663800 0.207500

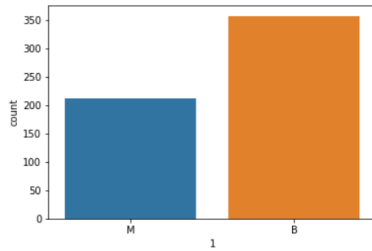
```

The **describe** function describes each column in the data and gives simple statistics. There are only 2 unique classes in target variable, which is in agreement with the data description. Since we do not have column descriptions, it is difficult to comment on the distribution of the predictor variables.

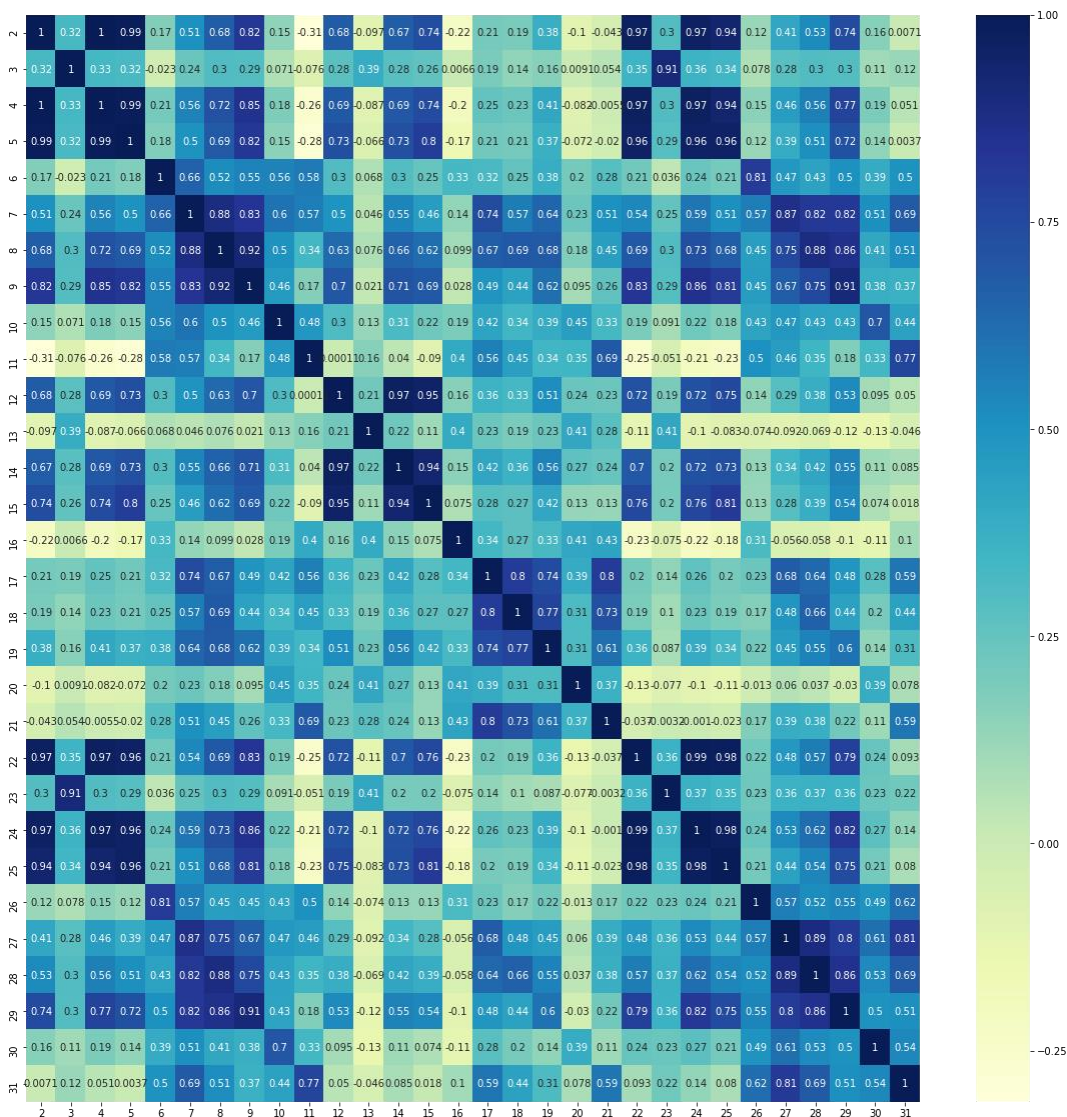
## 2) Exploratory data analysis & Visualization

```
import matplotlib.pyplot as plt
#pyplot is matplotlib's plotting framework
import seaborn as sns
# Seaborn is a Python data visualization library based on matplotlib.

sns.countplot(df[1])
plt.show()
```



Looking at a histogram of the target variable, we observe that there are 357(63%) instances of benign and 212(37%) instances of malignancy, indicating slight a class imbalance in the data. While this is not an ideal 50:50 distribution, we continue without making any changes to the distribution. There are certain techniques like up-sampling and down-sampling that can be used to alter the class distribution of target variable, but for the sake of simplicity, we leave the target variable as is.



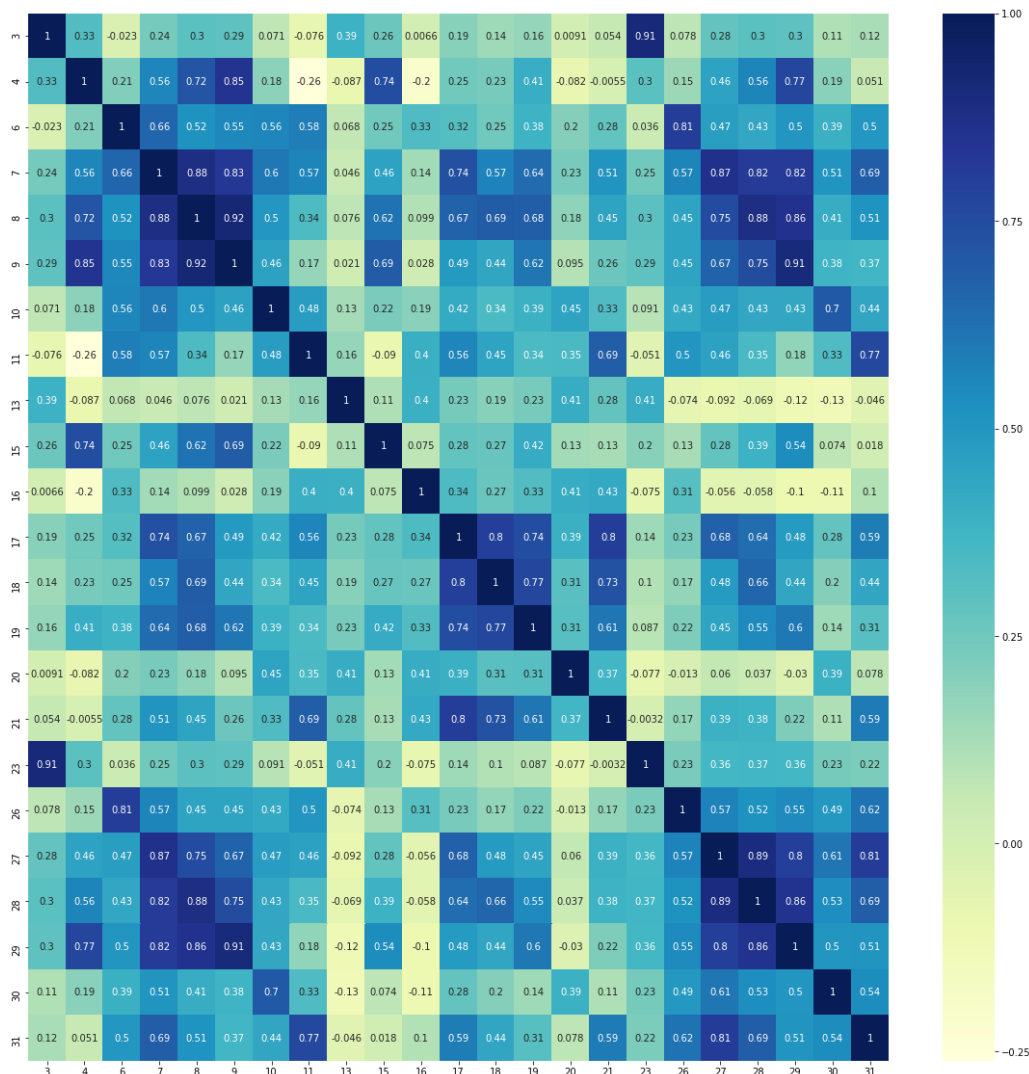


Looking at the correlation heat map, we observe possible cases of multicollinearity. This can be seen in the rather high correlation values(>0.93) between the following pairs of variables – (2,5) (2,22) (2,24) (2,25) (4,5) (4,22) (4,24) (4,25) (5,22) (5,24) (5,25) (12,14) (12,15) (14,15) (22,24) (22,25).

Since multicollinearity may lead to overfitting, we drop the following variables that seem to be causing the multicollinearity – 2,5,12,14,22,24,25

```
df2= df.copy()
df2.drop(df2.columns[[2,5,12,14,22,24,25]],axis=1, inplace = True)
```

The correlation heatmap after dropping the variables is below, we notice that the extreme correlations of >0.93 are removed, which should have reduced the problem of multicollinearity:



### 3) Splitting data and standardization

In this step, the data is split into train and test in the ratio 70:30 respectively. The stratified sampling is used to make sure the distribution of both classes is consistent across train and test data. Since the KNN modeling uses distance as a measure to find nearest neighbors, it is important to standardize the variables so that the distance is unaffected by scaling of the variables. We use the Z score scaling method to standardize the variables.

## Splitting data and standardization

```
#split dataset in features and target variable
X = df2.iloc[:,2:]# Features
y = df2[1] # Target variable
classes = ['M', 'B' ]
##### Split the Data #####

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=1, stratify=y)

##### Normalization #####

from sklearn.preprocessing import StandardScaler # Standardize features by removing the mean and scaling to unit variance

sc = StandardScaler()
sc.fit(X_train) # Compute the mean and std to be used for later scaling.

X_train_std = sc.transform(X_train) # Perform standardization of train set X by centering and scaling
X_test_std = sc.transform(X_test) # Perform standardization of test set X by centering and scaling
```

## 4) KNN model induction

```
##### Train the Model #####

from sklearn import neighbors, datasets

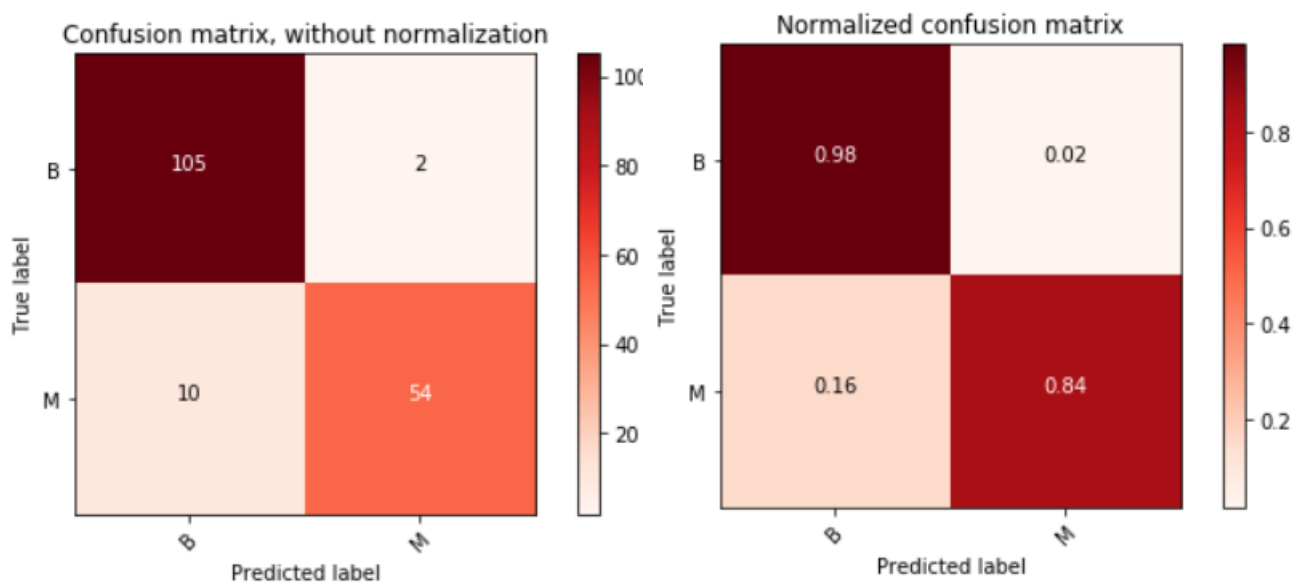
# KNeighborsClassifier is a classifier implementing the k-nearest neighbors vote.
# Learn more about it here https://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html

# Set parameters of KNeighborsClassifier
knn = neighbors.KNeighborsClassifier(n_neighbors=3, #n_neighbors is the k in the kNN
                                   p=2,
                                   metric='minkowski', weights='distance') #The default metric is minkowski, which is a generalization
                                   # with p=2 is equivalent to the standard Euclidean distance.
                                   # with p=1 is equivalent to the Mahattan distance.

# Train the model
knn = knn.fit(X_train_std, y_train)
```

In this step, we train a KNN classifier with  $k=3$  on the standardized training and test data. The distance metric used is a generalization of the Euclidean distance, which is specified by  $p=2$ . The weights parameter is specified as “distance”, which uses a weighted voting for nearest neighbors, indicating that points nearer will have a greater contribution in class prediction. All other modeling parameters are left as default.

## 5) KNN model evaluation



```

Accuracy (out-of-sample): 0.93
Accuracy (in-sample): 1.00
      precision    recall  f1-score   support

     B         0.91         0.98         0.95         107
     M         0.96         0.84         0.90          64

```

The KNN model has an out of sample accuracy score of 93%. It does a pretty good job with the classification task on both classes, with an F measure of 0.95 for B class and 0.9 for M class, which are very high values, indicating the model is performing well. The precision measure is very high (>0.91) for both classes. The recall metric is nearly =1 for B class (0.98) but not as good for M class (0.84). Looking deeper into this, we see from the confusion matrix that the model wrongly classifies 10 cases of malignancy as Benign. This is an issue since we need to consider the cost of misclassification in both classes. We will look to see if logistic regression can perform better in classifying more malignant classes accurately.

## 6) Logistic Regression model induction and evaluation

In this step, we train a logistic regression model to carry out the same classification process. It is important to note that we do not use the standardized form of the data here, but rather use the raw unstandardized data. In this case, we use the default penalty option of L2 norm. The C value used is 10000, which is the value for inverse of regularization strength. All other parameters in the modeling are left as default.

```

from sklearn import linear_model
clf = linear_model.LogisticRegression(C=1e5)

clf = clf.fit(X_train, y_train)
print('The weights of the attributes are:', clf.coef_)
##### Apply the Logistic Regression Model #####

y_pred = clf.predict(X_test) # Classification prediction
y_pred_prob = clf.predict_proba(X_test) # Class probabilities
print(y_pred[0], y_pred_prob[0], np.sum(y_pred_prob[0]))

##### Evaluate the Logistic Regression Model #####

# Build a text report showing the main classification metrics (out-of-sample performance)
print(classification_report(y_test, y_pred, target_names=classes))

# Compute confusion matrix to evaluate the accuracy of a classification
cnf_matrix = confusion_matrix(y_test, y_pred)
np.set_printoptions(precision=2)

#cnf_matrix = confusion_matrix(y_test, y_pred)

# Plot non-normalized confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=classes,
                      title='Confusion matrix, without normalization')

# Plot normalized confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=classes, normalize=True,
                      title='Normalized confusion matrix')

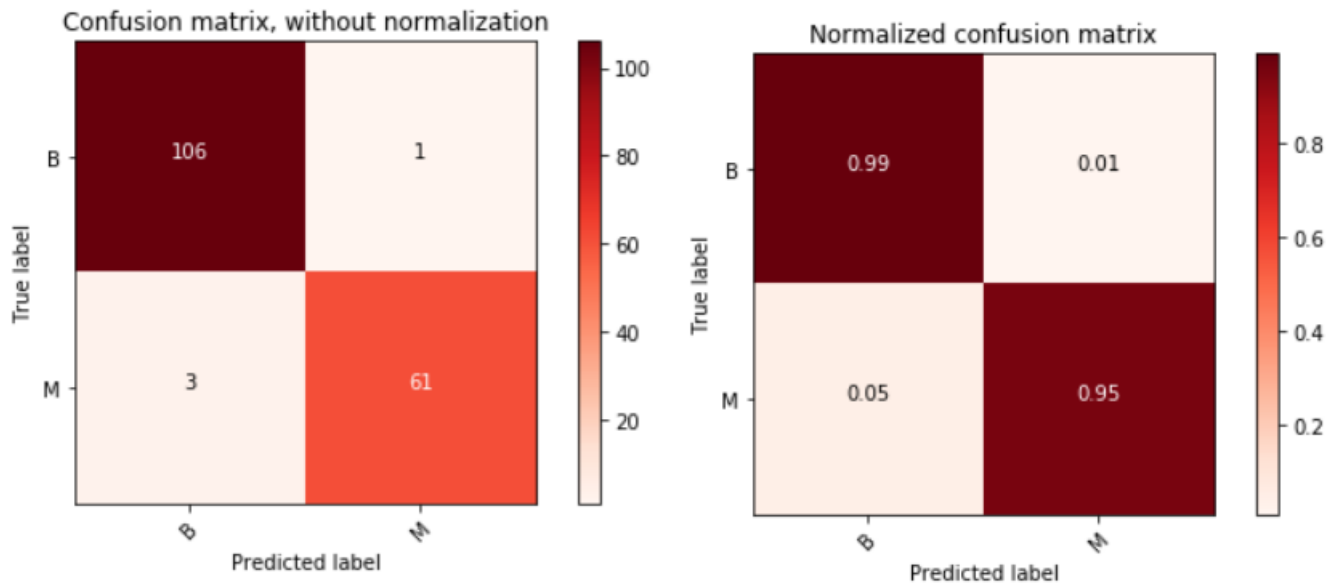
plt.show()

```

```

      precision    recall  f1-score
Accuracy (out-of-sample): 0.98  B         0.97         0.99         0.98
Accuracy (in-sample): 1.00    M         0.98         0.95         0.97

```



The logistic regression model has an out of sample accuracy of 98%. The F measure for B class is 0.98, and for M class is 0.97. These measures are almost =1, indicating the model is performing very well in classifying both classes. Unlike the KNN model, where recall for class M was lower than that of B, here recall scores for both classes are very high (0.95 for M and 0.99 for B). The precision scores for both classes are near perfect as well (0.97 for B and 0.98 for M). On all metrics, logistic regression model has a superior performance in comparison to KNN model.

#### 7) KNN vs Logistic regression

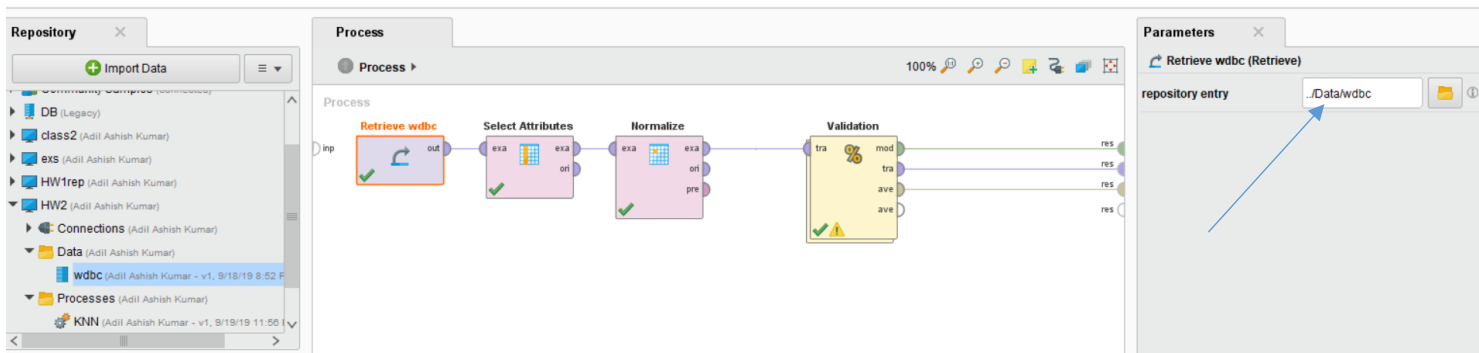
In terms of out of sample accuracy, Logistic regression (98%) performs better than KNN (93%). The F measures for Logistic regression for both classes “B”, “M” are better than KNN. Just on performance, it seems like logistic regression performs better than KNN.

An important differentiating factor is the F measure for both models. In KNN, the model predicts about 16% of the malignant cases as benign. This is a major issue, as there are serious costs and consequences involved in wrongly predicting occurrence of a disease like cancer. Thus, I would prefer to use Logistic regression model, as it has much better accuracy in predicting both benign and malignant classes, which are evident from the F measures for B and M classes.

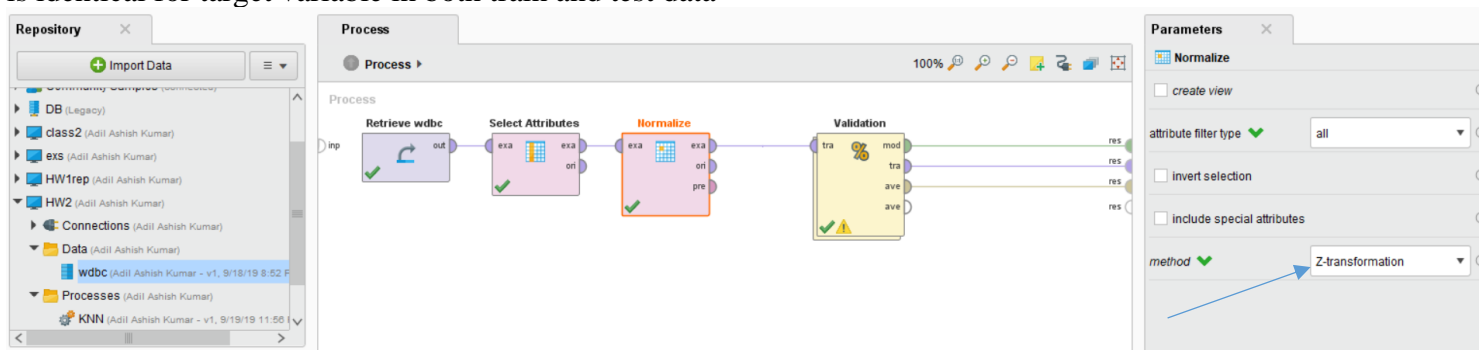
An important consideration is that logistic regression may not be very suitable for very large datasets, so the size of data should be considered when choosing the model for deployment. Another important point is that logistic regression model may not be able to handle complexities in data like missing values, non linearities etc, so additional data preparation maybe needed if the data has such issues. Model comprehensibility is also an important factor and logistic regression is not an easy model to understand. Here I assume that the dataset size is small enough to be suitable for logistic regression. I also assume that model generalization performance is more important than model interpretability. Thus, I would suggest using the logistic regression model.

#### Rapid miner implementation

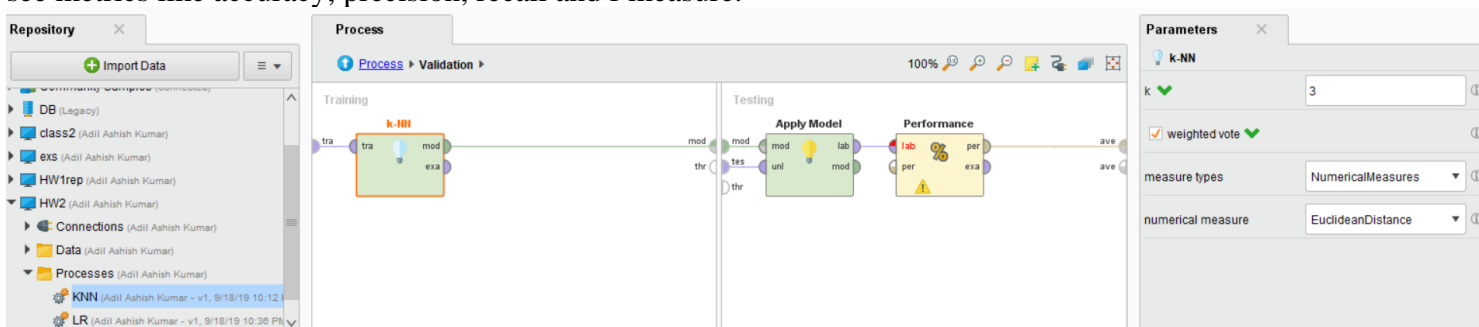
##### 1) KNN



The modeling process in rapidminer is similar to that of python. First the data is read in using the retrieve module. The select attribute module is used to drop the case ID and other columns that were causing multicollinearity with help of subset option. Then the data is normalized using the Z-score method, where mean of data is subtracted from all values and divided by the std deviation. We use a 70:30 split for train and test data respectively in the validation module. The stratified sampling method is used so that class distribution is identical for target variable in both train and test data



Under the validation module, we use the knn module, the apply module and the performance module. In KNN module, we use K=3, choose the distance measure as Euclidean distance. Since all predictor attributes in our model are numerical, we choose the measure type as “numerical”. In the performance module, I have chosen to see metrics like accuracy, precision, recall and f measure.



On running the model, we get an out of sample accuracy of 94.15%. The F measure for class “M” is 91.8%. The F measure and accuracy are high (>0.9), indicating the model is doing a good job with the classification task. The confusion matrix shows that the model is doing a good job with predicting both classes.



Views: DesignResultsTurbo PrepAuto Model

Find data, open

Result History

PerformanceVector (Performance)

ExampleSet (Select Attributes)

LogisticRegression (Logistic Regression)

Performance

Description

Annotations

Criterion

accuracy

classification error

precision

recall

f measure

Table View

Plot View

accuracy: 95.32%

	true M	true B	class precision
pred. M	61	5	92.42%
pred. B	3	102	97.14%
class recall	95.31%	95.33%	

Result History

PerformanceVector (Performance)

ExampleSet (Select Attributes)

LogisticRegression (Logistic Regression)

Performance

Description

Annotations

Criterion

accuracy

classification error

precision

recall

f measure

Table View

Plot View

f\_measure: 93.85% (positive class: M)

	true B	true M	class precision
pred. B	102	3	97.14%
pred. M	5	61	92.42%
class recall	95.33%	95.31%	

The accuracy out of sample is 95.32%. The F measure for class M is 93.85%. In this case as well, the performance of the logistic regression model is better than the KNN model on both out of sample accuracy and F measures for both classes.